

**15-441**

# **Network Programming**

## **September 6, 2005**

**David Murray**

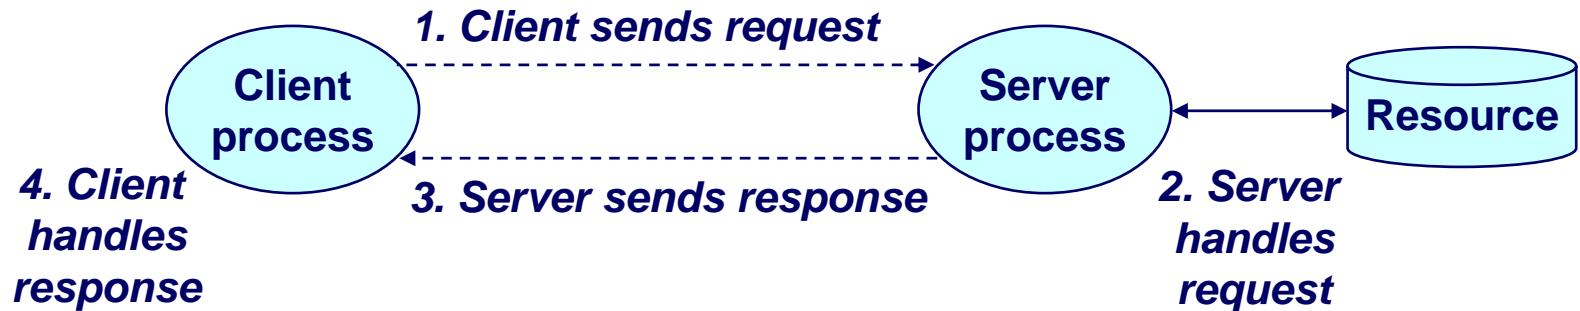
**Slides based on those of Dave Maltz,  
Randy Bryant, Geoff Langale,  
and the 15-213 crew**

### **Topics**

- **Programmer's view of the Internet**
- **Sockets interface**
- **Writing clients and servers**
- **Concurrency with I/O multiplexing**
- **Debugging With GDB**
- **Version Control (RCS/CVS)**
- **Tips from the trenches: Projects 1 & 2**

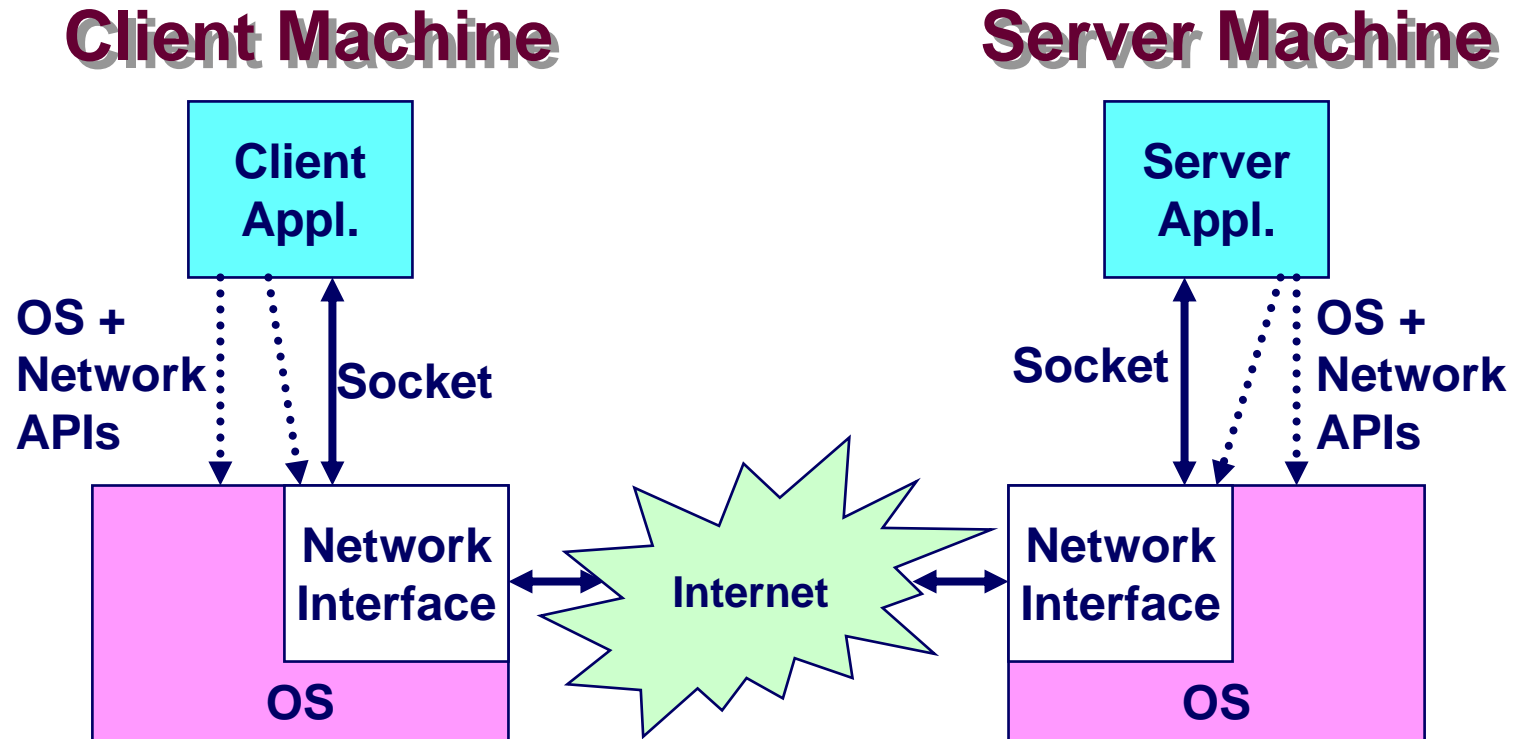
# A Client-Server Exchange

- A **server** process and one or more **client** processes
- Server manages some **resource**.
- Server provides **service** by manipulating resource for clients.



*Note: clients and servers are processes running on hosts (can be the same or different hosts).*

# Network Applications



## Access to Network via Program Interface


- Sockets make network I/O look like files
- Call system functions to control and communicate
- Network code handles issues of routing, segmentation.

# Internet Connections (TCP/IP)

Two common paradigms for clients and servers communication

- Datagrams (UDP protocol SOCK\_DGRAM)
- Connections (TCP protocol, SOCK\_STREAM)

Connections are point-to-point, full-duplex (2-way communication), and reliable. (TODAY'S TOPIC!)



*Note: 3479 is an ephemeral port allocated by the kernel*

*Note: 80 is a well-known port associated with Web servers*  
15-441, Spring 2005

# Clients

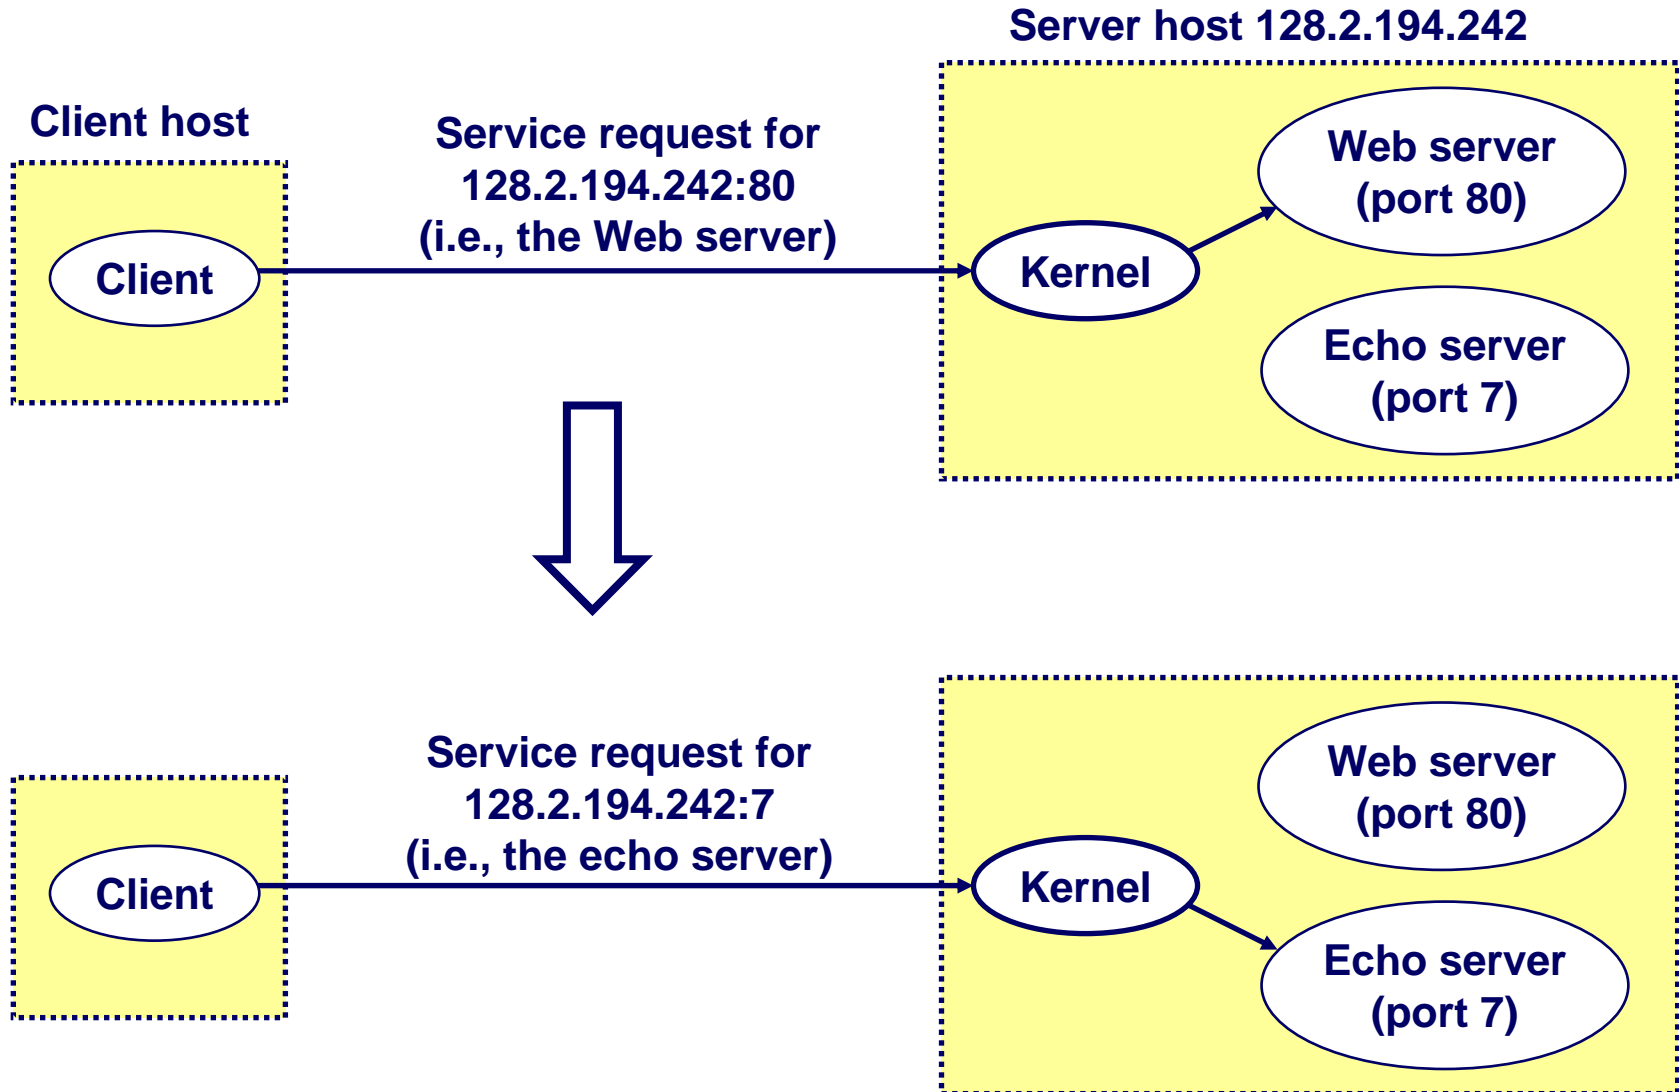
## Examples of client programs

- Web browsers, ftp, telnet, ssh

## How does a client find the server?

- The IP address in the server socket address identifies the host *(more precisely, an adaptor on the host)*
- The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
- Examples of well known ports
  - Port 7: Echo server
  - Port 23: Telnet server
  - Port 25: Mail server
  - Port 80: Web server

# Using Ports to Identify Services



# Servers

**Servers are long-running processes (daemons).**

- Created at boot-time (typically) by the init process (process 1)
- Run continuously until the machine is turned off.

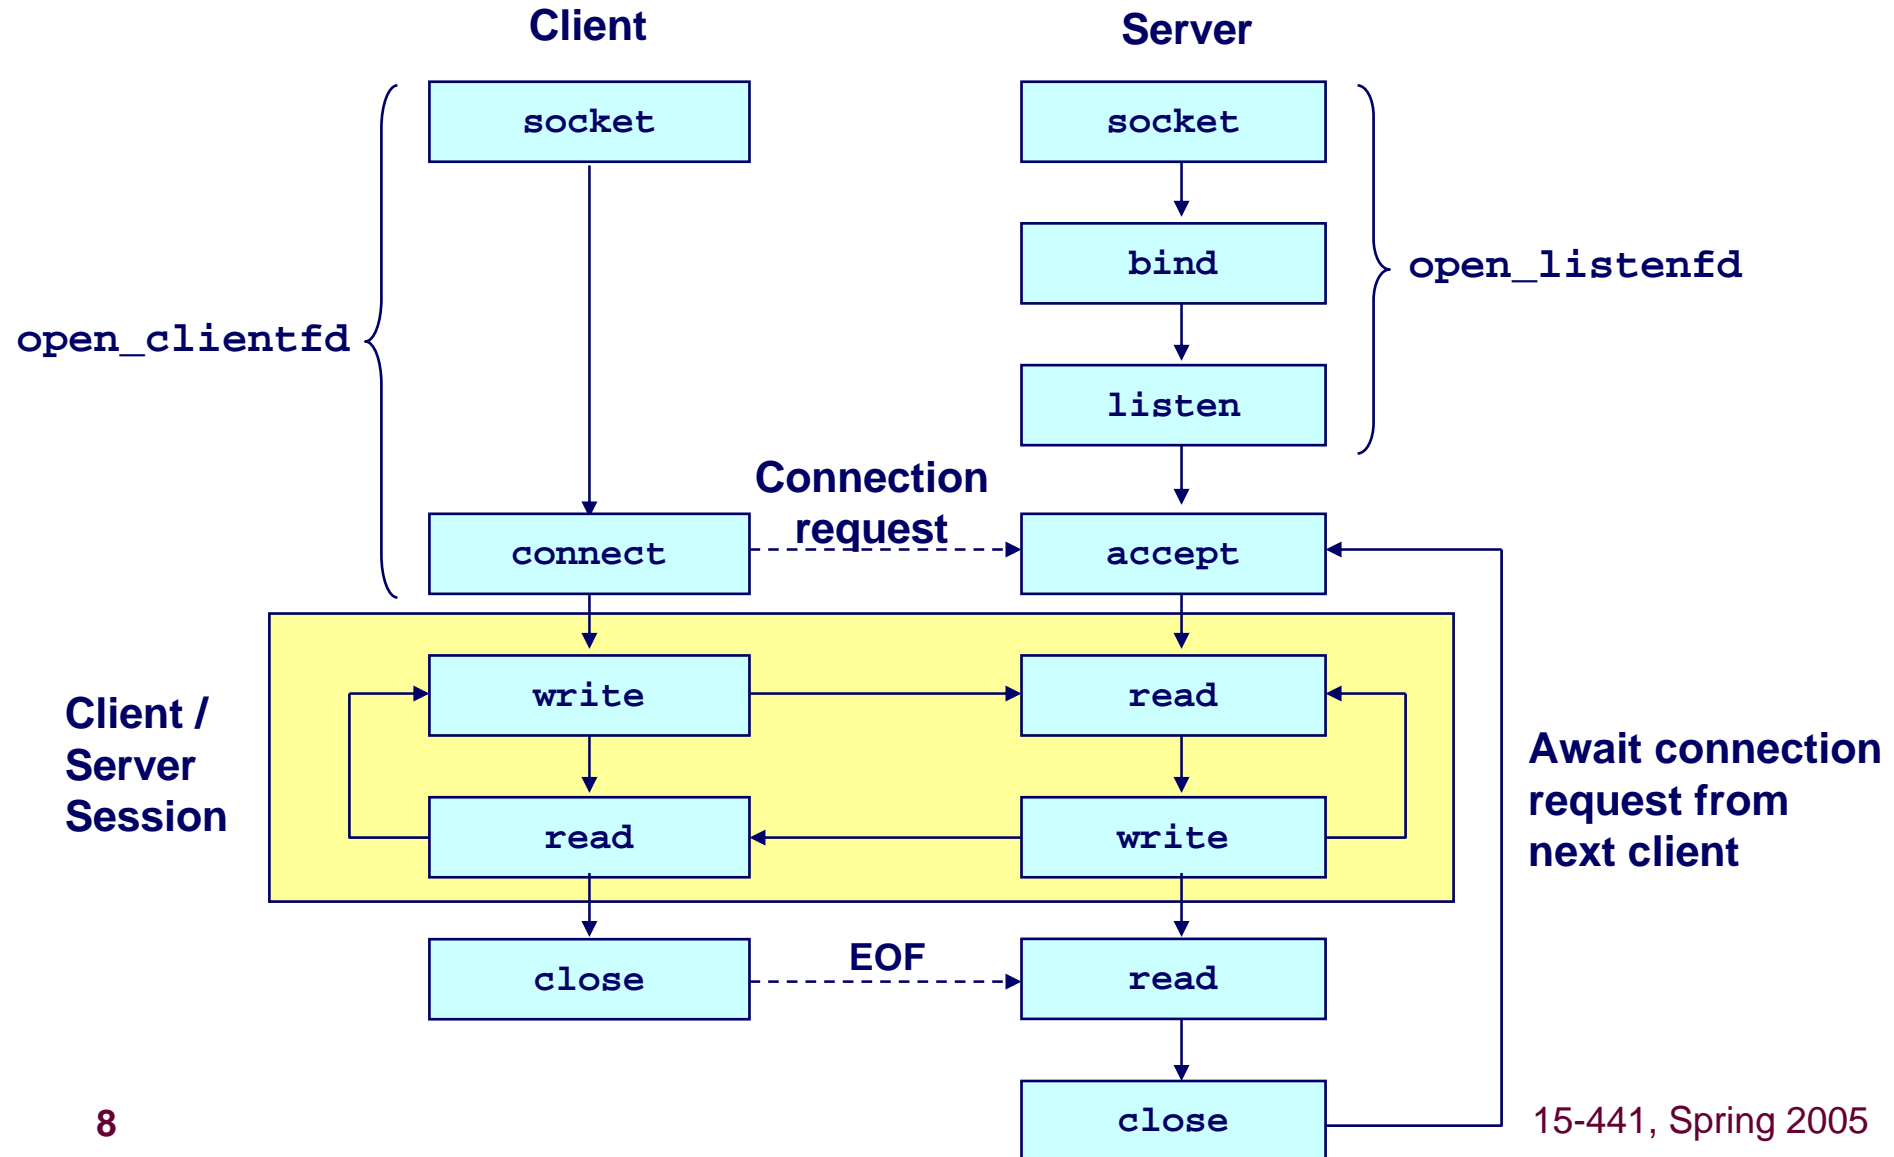
**Each server waits for requests to arrive on a well-known port associated with a particular service.**

- Port 7: echo server
- Port 23: telnet server
- Port 25: mail server
- Port 80: HTTP server

See `/etc/services` for a comprehensive list of the services available on a Linux machine.

**A machine that runs a server process is also often referred to as a “server.”**

# Overview of the Sockets Interface





# Sockets

## What is a socket?

- To the kernel, a socket is an endpoint of communication.
- To an application, a socket is a file descriptor that lets the application read/write from/to the network.
  - Remember: All Unix I/O devices, including networks, are modeled as files.

**Clients and servers communicate with each by reading from and writing to socket descriptors.**

**The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.**

# Socket Programming Cliches

## Network Byte Ordering

- Network is big-endian, host may be big- or little-endian
- Functions work on 16-bit (short) and 32-bit (long) values
- htons() / htonl() : convert host byte order to network byte order
- ntohs() / ntohl(): convert network byte order to host byte order
- Use these to convert network addresses, ports, ...

## Structure Casts

- You will see a lot of 'structure casts'

```
struct sockaddr_in serveraddr;  
/* fill in serveraddr with an address */  
...  
/* Connect takes (struct sockaddr *) as its second argument */  
connect(clientfd, (struct sockaddr *) &serveraddr,  
        sizeof(serveraddr));  
...
```

# Socket Programming Help

**man is your friend (aka RTFM)**

- **man accept**
- **man select**
- **Etc.**

**The manual page will tell you:**

- **What #include<> directives you need at the top of your source code**
- **The type of each argument**
- **The possible return values**
- **The possible errors (in errno)**

# Socket Address Structures

## Generic socket address:

- For address arguments to connect, bind, and accept.

```
struct sockaddr {  
    unsigned short  sa_family;    /* protocol family */  
    char            sa_data[14]; /* address data.  */  
};
```

## Internet-specific socket address:

- Must cast (sockaddr\_in \*) to (sockaddr \*) for connect, bind, and accept.

```
struct sockaddr_in {  
    unsigned short  sin_family; /* address family (always AF_INET) */  
    unsigned short  sin_port;   /* port num in network byte order */  
    struct in_addr  sin_addr;   /* IP addr in network byte order */  
    unsigned char   sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```

# Reliable I/O (RIO) Summary

## I/O Package Developed by David O'Hallaron

- <http://csapp.cs.cmu.edu/public/code.html> (csapp.{h,c})
- Allows mix of buffered and unbuffered I/O

## Important Functions

`rio_writen(int fd, void *buf, size_t n)`

- Writes `n` bytes from buffer `buf` to file `fd`.

`rio_readlineb(rio_t *rp, void *buf, size_t maxn)`

- Read complete text line from file `rp` into buffer `buf`.
  - » Line must be terminated by newline (`\n`) character
- Up to maximum of `maxn` bytes

## Used Here For Illustrative Purposes Only

- You may want to use `read()/write()` for your projects instead
- You will need to check error returns
- Reading a whole line won't always make sense (more later)
- **NOTE: RIO functions capitalize first letter!! You must fix this!**
  - `Accept()` .vs. `accept()`

# Echo Client Main Routine

```
#include "csapp.h"

/* usage: ./echoclient host port */
int main(int argc, char **argv)
{
    int clientfd, port;
    char *host, buf[MAXLINE];
    rio_t rio;

    host = argv[1];
    port = atoi(argv[2]);

    clientfd = Open_clientfd(host, port);
    Rio_readinitb(&rio, clientfd);

    while (Fgets(buf, MAXLINE, stdin) != NULL) {
        Rio_writen(clientfd, buf, strlen(buf));
        Rio_readlineb(&rio, buf, MAXLINE);
        Fputs(buf, stdout);
    }
    Close(clientfd);
    exit(0);
}
```

Send line to  
server



Receive line  
from server



# Client-side Programming

# Echo Client: open\_clientfd

```
int open_clientfd(char *hostname, int port)
{
    int clientfd;
    struct hostent *hp;
    struct sockaddr_in serveraddr;

    if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1; /* check errno for cause of error */

    /* Fill in the server's IP address and port */
    if ((hp = gethostbyname(hostname)) == NULL)
        return -2; /* check h_errno for cause of error */
    bzero((char *) &serveraddr, sizeof(serveraddr));
    serveraddr.sin_family = AF_INET;
    serveraddr.sin_port = htons(port);
    bcopy((char *)hp->h_addr,
          (char *)&serveraddr.sin_addr.s_addr, hp->h_length);

    /* Establish a connection with the server */
    if (connect(clientfd, (struct sockaddr *) &serveraddr,
                sizeof(serveraddr)) < 0)
        return -1;
    return clientfd;
}
```

This function opens a connection from the client to the server at hostname:port



# Echo Client: `open_clientfd` (`socket`)

**`socket` creates a socket descriptor on the client.**

- `AF_INET`: indicates that the socket is associated with Internet protocols.
- `SOCK_STREAM`: selects a reliable byte stream connection.

```
int clientfd; /* socket descriptor */  
  
if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)  
    return -1; /* check errno for cause of error */  
  
... (more)
```

# Echo Client: `open_clientfd` (`gethostbyname`)

The client then builds the server's Internet address.

```
int clientfd;                /* socket descriptor */
struct hostent *hp;          /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */

...

/* fill in the server's IP address and port */
if ((hp = gethostbyname(hostname)) == NULL)
    return -2; /* check h_errno for cause of error */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
bcopy((char *)hp->h_addr,
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
serveraddr.sin_port = htons(port);
```

# Echo Client: `open_clientfd` (`connect`)

Finally the client creates a connection with the server.

- Client process suspends (blocks) until the connection is created.
- After resuming, the client is ready to begin exchanging messages with the server via Unix I/O calls on descriptor `sockfd`.

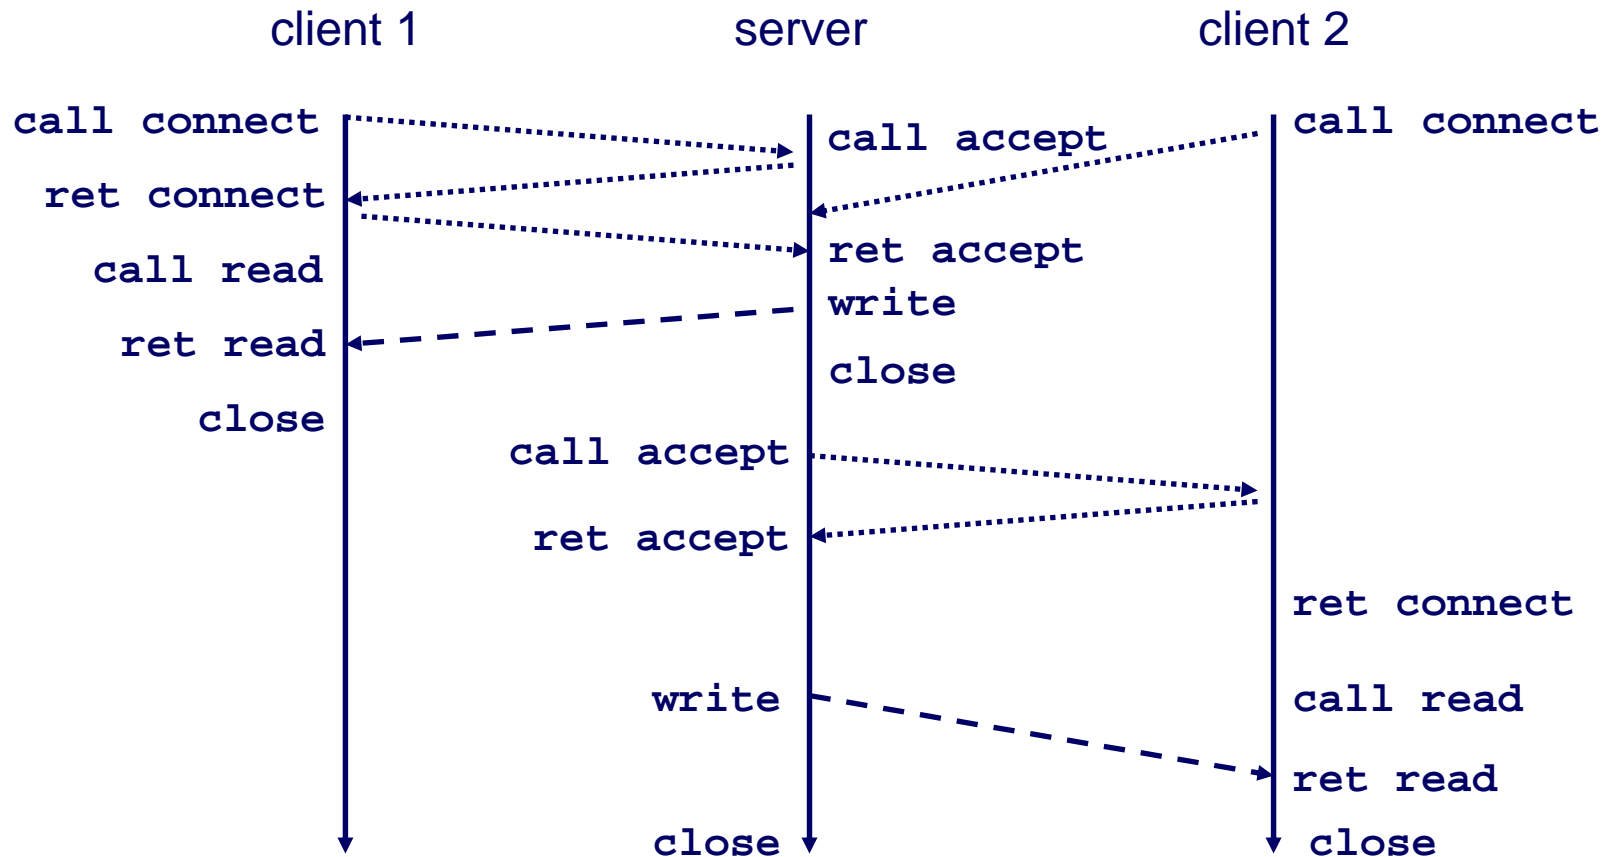
```
int clientfd;                /* socket descriptor */
struct sockaddr_in serveraddr; /* server address */
typedef struct sockaddr SA;   /* generic sockaddr */
...
/* Establish a connection with the server */
if (connect(clientfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
return clientfd;
}
```

# Server-side Programming

# Servers and sockets – 1 isn't enough

Server must be able to handle multiple requests

Where should pending connections be queued up?



# Connected vs. Listening Descriptors

## Listening descriptor

- End point for client connection requests.
- Created once and exists for lifetime of the server.

## Connected descriptor

- End point of the connection between client and server.
- A new descriptor is created each time the server accepts a connection request from a client.
- Exists only as long as it takes to service client.


## Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously.

# Echo Server: accept Illustrated



**1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`.**



**2. Client makes connection request by calling and blocking in `connect`.**



**3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`.**

# Echo Server: Main Loop

The server loops endlessly, waiting for connection requests, then reading input from the client, and echoing the input back to the client.

```
main() {  
  
    /* create and configure the listening socket */  
  
    while(1) {  
        /* Accept(): wait for a connection request */  
        /* echo(): read and echo input lines from client til EOF */  
        /* Close(): close the connection */  
    }  
}
```



# Echo Server: open\_listenfd

```
int open_listenfd(int port)
{
    int listenfd, optval=1;
    struct sockaddr_in serveraddr;

    /* Create a socket descriptor */
    if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1;

    /* Eliminates "Address already in use" error from bind. */
    if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
                    (const void *)&optval , sizeof(int)) < 0)
        return -1;

    ... (more)
```

# Echo Server: open\_listenfd (cont)

```
...

/* Listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
serveraddr.sin_port = htons((unsigned short)port);
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;

/* Make it a listening socket ready to accept
   connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;

return listenfd;
}
```

# Echo Server: `open_listenfd` (`socket`)

**`socket` creates a socket descriptor on the server.**

- `AF_INET`: indicates that the socket is associated with Internet protocols.
- `SOCK_STREAM`: selects a reliable (TCP) byte stream connection.

```
int listenfd; /* listening socket descriptor */  
  
/* Create a socket descriptor */  
if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)  
    return -1;
```

# Echo Server: `open_listenfd` (initialize socket address)

Next, we initialize the socket with the server's Internet address (IP address and port)

```
struct sockaddr_in serveraddr; /* server's socket addr */
...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
serveraddr.sin_port = htons((unsigned short)port);
```

**IP addr and port stored in network (big-endian) byte order**

- `htonl()` converts longs from host byte order to network byte order.
- `htons()` converts shorts from host byte order to network byte order.

# Echo Server: `open_listenfd` (`bind`)

`bind` associates the socket with the socket address we just created.

```
int listenfd;                /* listening socket */
struct sockaddr_in serveraddr; /* server's socket addr */

...
/* listenfd will be an endpoint for all requests to port
   on any IP address for this host */
if (bind(listenfd, (struct sockaddr *)&serveraddr,
          sizeof(serveraddr)) < 0)
    return -1;
```

# Echo Server: `open_listenfd` (`listen`)

`listen` indicates that this socket will accept connection (`connect`) requests from clients.

```
int listenfd; /* listening socket */

...
/* Make it a listening socket ready to accept connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;
return listenfd;
}
```

We're finally ready to enter the main server loop that accepts and processes client connection requests.

# Echo Server: accept

**accept ( ) blocks waiting for a connection request.**

```
int listenfd; /* listening descriptor */
int connfd;   /* connected descriptor */
struct sockaddr_in clientaddr;
int clientlen;

clientlen = sizeof(clientaddr);
connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
```

**accept returns a *connected descriptor* (connfd) with the same properties as the *listening descriptor* (listenfd)**

- Returns when the connection between client and server is created and ready for I/O transfers.
- All I/O with the client will be done via the connected socket.

**accept also fills in client's IP address.**

# Echo Server: Main Routine

```
int main(int argc, char **argv) {
    int listenfd, connfd, port, clientlen;
    struct sockaddr_in clientaddr;
    struct hostent *hp;
    char *haddrp;

    port = atoi(argv[1]); /* the server listens on a port passed
                           on the command line */
    listenfd = open_listenfd(port);

    while (1) {
        clientlen = sizeof(clientaddr);
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                           sizeof(clientaddr.sin_addr.s_addr), AF_INET);
        haddrp = inet_ntoa(clientaddr.sin_addr);
        printf("Fd %d connected to %s (%s:%s)\n",
               connfd, hp->h_name, haddrp, ntohs(clientaddr.sin_port));
        echo(connfd);
        Close(connfd);
    }
}
```



# Echo Server: Identifying the Client

The server can determine the domain name, IP address, and port of the client.

```
struct hostent *hp; /* pointer to DNS host entry */
char *haddrp;      /* pointer to dotted decimal string */

hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                  sizeof(clientaddr.sin_addr.s_addr), AF_INET);
haddrp = inet_ntoa(clientaddr.sin_addr);
printf("Fd %d connected to %s (%s:%s)\n",
       connfd, hp->h_name, haddrp, ntohs(clientaddr.sin_port));
```

# Echo Server: echo

The server uses RIO to read and echo text lines until EOF (end-of-file) is encountered.

- EOF notification caused by client calling `close(clientfd)`.
- **IMPORTANT:** EOF is a condition, not a particular data byte.


```
void echo(int connfd)
{
    size_t n;
    char buf[MAXLINE];
    rio_t rio;

    Rio_readinitb(&rio, connfd);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        printf("server received %d bytes\n", n);
        Rio_writen(connfd, buf, n);
    }
}
```

Receive line  
from client



Send line to  
client



# Running Echo Client/Server

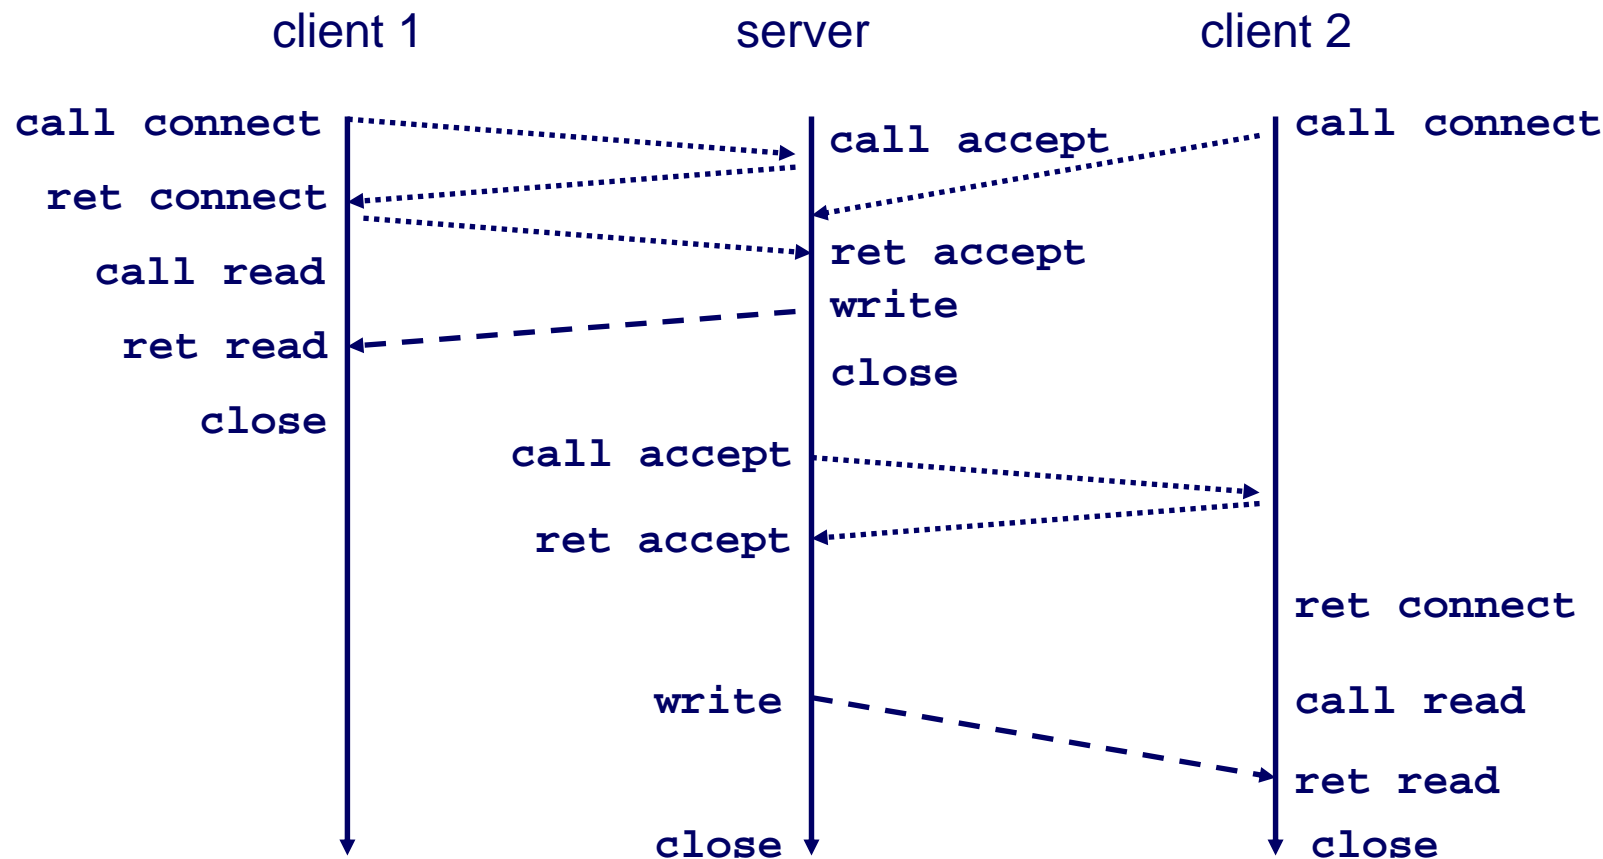
```
[bryant@bryant echo]$ ./echoservers 15441  
fd 4 connected to BRYANT-TP2.VLSI.CS.CMU.EDU (128.2.222.198:3507)  
Server received 12 (12 total) bytes on fd 4
```

```
[bryant@bryant-tp2 echo]$ ./echoclient bryant.vlsi.cs.cmu.edu 15441  
hello world  
hello world
```

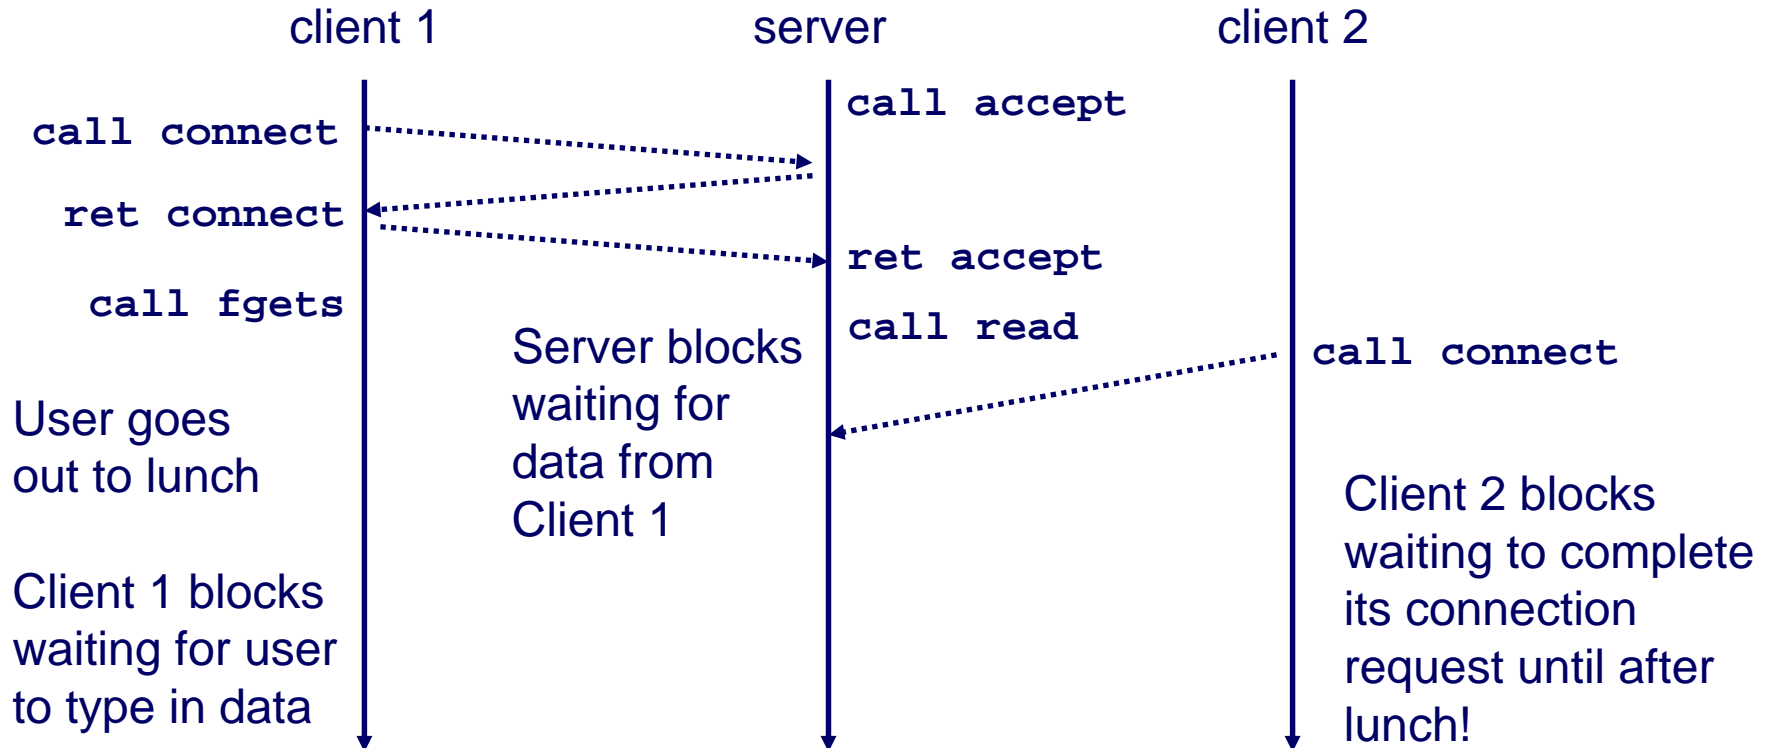
# Types of Server Implementations

# Iterative Servers

Iterative servers process one request at a time.



# Fundamental Flaw of Iterative Servers

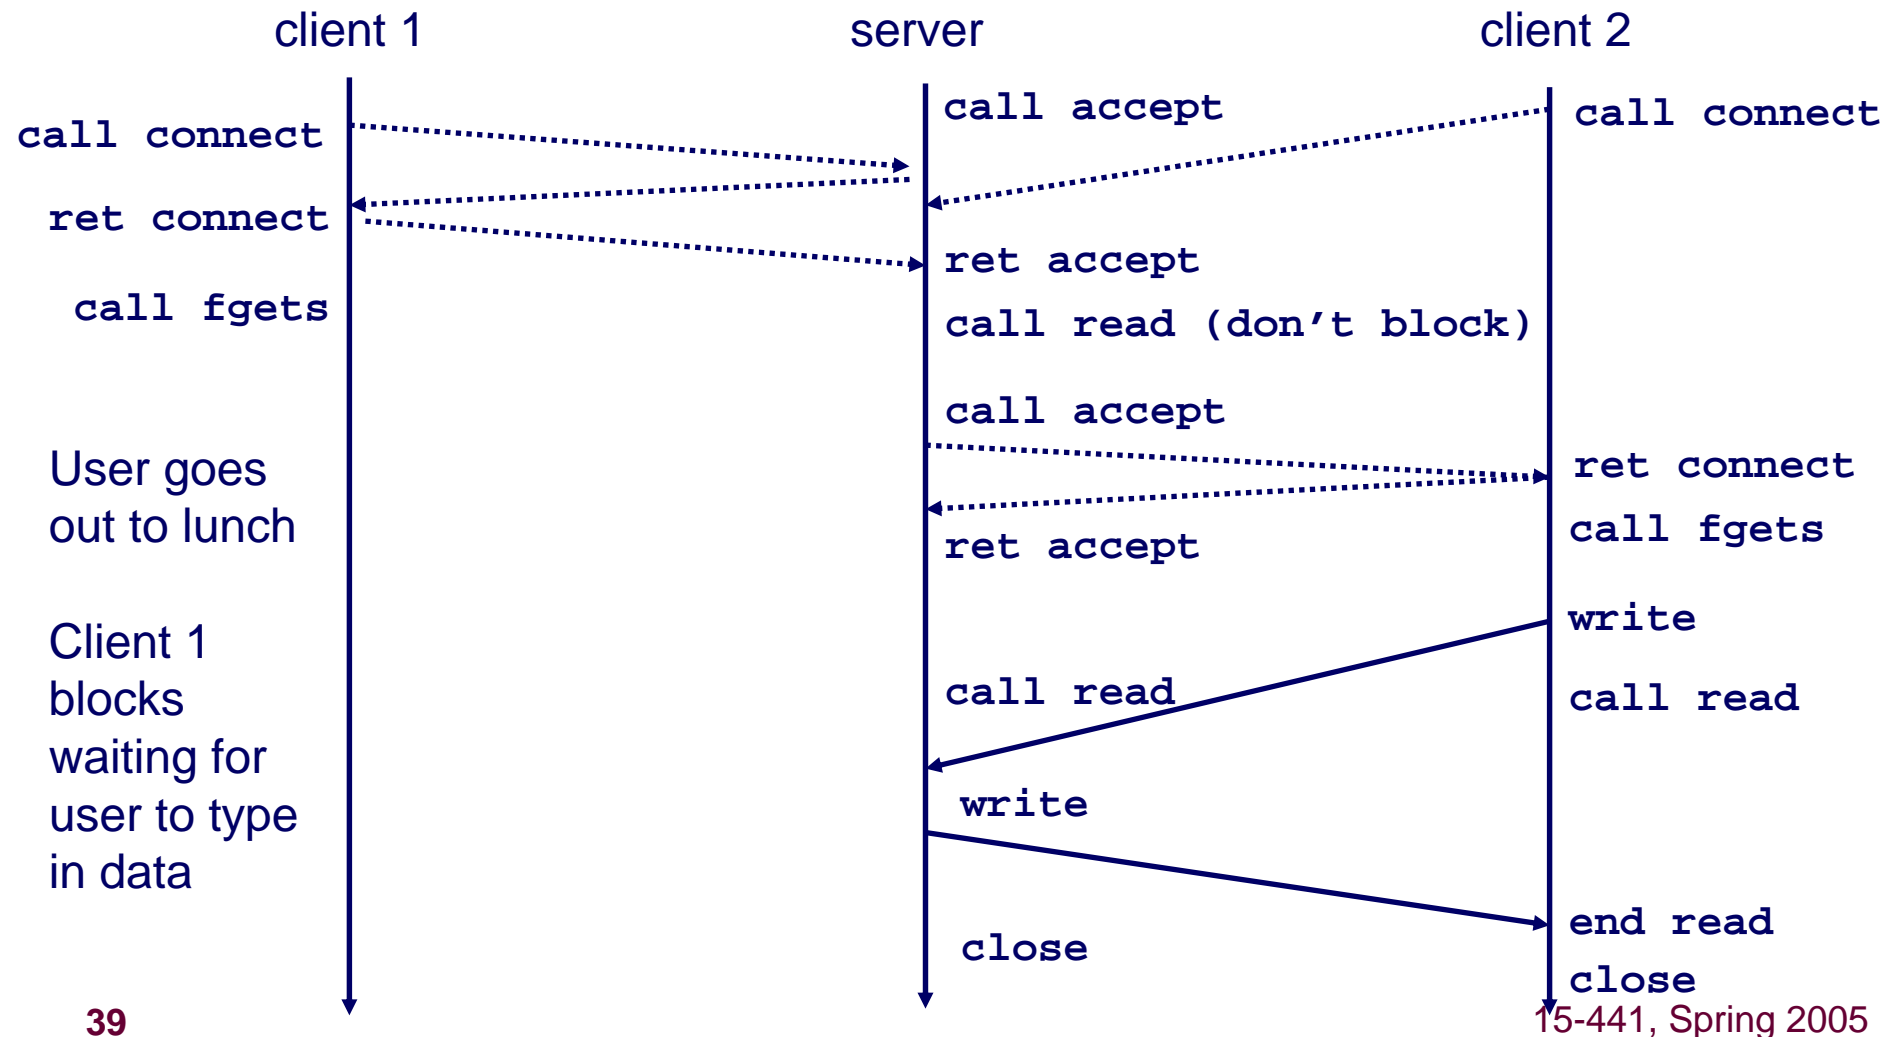


**Solution: use *concurrent servers* instead.**

- Concurrent servers use multiple concurrent flows to serve multiple clients at the same time.

# Concurrent Servers

**Concurrent servers handle multiple requests concurrently.**



# Possible Mechanisms for Creating Concurrent Flows

## 1. Processes

- Kernel automatically interleaves multiple logical flows.
- Each flow has its own private address space.

## 2. I/O multiplexing with `select ( )`

**Our Focus**

- User manually interleaves multiple logical flows.
- Each flow shares the same address space.
- Popular for high-performance server designs.

## 3. Threads

- Kernel automatically interleaves multiple logical flows.
- Each flow shares the same address space.



# Event-Based Concurrent Servers Using I/O Multiplexing

**Maintain a pool of connected descriptors.**

**Repeat the following forever:**

- Use the Unix `select` function to block until:
  - (a) New connection request arrives on the listening descriptor.
  - (b) New data arrives on an existing connected descriptor.
- If (a), add the new connection to the pool of connections.
- If (b), read any available data from the connection
  - Close connection on EOF and remove it from the pool.

# The select Function

`select()` sleeps until one or more file descriptors in the set `readset` ready for reading or one or more descriptors in `writeset` ready for writing

```
#include <sys/select.h>
```

```
int select(int maxfdp1, fd_set *readset, fd_set *writeset,  
NULL, NULL);
```

## `readset`

- Opaque bit vector (max `FD_SETSIZE` bits) that indicates membership in a *descriptor set*.
  - On Linux machines, `FD_SETSIZE` = 1024
- If bit `k` is 1, then descriptor `k` is a member of the descriptor set.
- When call `select`, should have `readset` indicate which descriptors to test

## `writeset`

- `writeset` is similar but refers to descriptors ready for writing

## `maxfdp1`

- Maximum descriptor in descriptor set plus 1.
- Tests descriptors 0, 1, 2, ..., `maxfdp1` - 1 for set membership.

`select()` returns the number of ready descriptors and keeps on each bit

42 of `readset` for which corresponding descriptor is ready

15-441, Spring 2005

# Macros for Manipulating Set Descriptors

```
void FD_ZERO(fd_set *fdset);
```

- Turn off all bits in fdset.

```
void FD_SET(int fd, fd_set *fdset);
```

- Turn on bit fd in fdset.

```
void FD_CLR(int fd, fd_set *fdset);
```

- Turn off bit fd in fdset.

```
int FD_ISSET(int fd, *fdset);
```

- Is bit fd in fdset turned on?

# Event-based Concurrent Echo Server

```
/*
 * echoservers.c - A concurrent echo server based on select
 */
#include "csapp.h"

typedef struct { /* represents a pool of connected descriptors */
    int maxfd;          /* largest descriptor in read_set */
    fd_set read_set;    /* set of all active descriptors */
    fd_set ready_set;   /* subset of descriptors ready for reading */
    int nready;         /* number of ready descriptors from select */
    int maxi;           /* highwater index into client array */
    int clientfd[FD_SETSIZE]; /* set of active descriptors */
    rio_t clientrio[FD_SETSIZE]; /* set of active read buffers */
} pool;

int byte_cnt = 0; /* counts total bytes received by server */
```

# Event-based Concurrent Server (cont)

```
int main(int argc, char **argv)
{
    int listenfd, connfd, clientlen = sizeof(struct sockaddr_in);
    struct sockaddr_in clientaddr;
    static pool pool;

    listenfd = Open_listenfd(argv[1]);
    init_pool(listenfd, &pool);

    while (1) {
        pool.ready_set = pool.read_set;
        pool.nready = Select(pool.maxfd+1, &pool.ready_set,
                             NULL, NULL, NULL);

        if (FD_ISSET(listenfd, &pool.ready_set)) {
            connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
            add_client(connfd, &pool);
        }
        check_clients(&pool);
    }
}
```

# Event-based Concurrent Server (cont)

```
/* initialize the descriptor pool */
void init_pool(int listenfd, pool *p)
{
    /* Initially, there are no connected descriptors */
    int i;
    p->maxi = -1;
    for (i=0; i< FD_SETSIZE; i++)
        p->clientfd[i] = -1;

    /* Initially, listenfd is only member of select read set */
    p->maxfd = listenfd;
    FD_ZERO(&p->read_set);
    FD_SET(listenfd, &p->read_set);
}
```

# Event-based Concurrent Server (cont)

```
void add_client(int connfd, pool *p) /* add connfd to pool p */
{
    int i;
    p->nready--;

    for (i = 0; i < FD_SETSIZE; i++) /* Find available slot */
        if (p->clientfd[i] < 0) {
            p->clientfd[i] = connfd;
            Rio_readinitb(&p->clientrio[i], connfd);

            FD_SET(connfd, &p->read_set); /* Add desc to read set */

            if (connfd > p->maxfd) /* Update max descriptor num */
                p->maxfd = connfd;
            if (i > p->maxi) /* Update pool high water mark */
                p->maxi = i;
            break;
        }
    if (i == FD_SETSIZE) /* Couldn't find an empty slot */
        app_error("add_client error: Too many clients");
}
```

# Event-based Concurrent Server (cont)

```
void check_clients(pool *p) { /* echo line from ready descs in pool p */
    int i, connfd, n;
    char buf[MAXLINE];
    rio_t rio;

    for (i = 0; (i <= p->maxi) && (p->nready > 0); i++) {
        connfd = p->clientfd[i];
        rio = p->clientrio[i];

        /* If the descriptor is ready, echo a text line from it */
        if ((connfd > 0) && (FD_ISSET(connfd, &p->ready_set))) {
            p->nready--;
            if ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
                byte_cnt += n;
                Rio_writen(connfd, buf, n);
            }
            else { /* EOF detected, remove descriptor from pool */
                Close(connfd);
                FD_CLR(connfd, &p->read_set);
                p->clientfd[i] = -1;
            }
        }
    }
}
```



# Pro and Cons of Event-Based Designs

- + One logical control flow.
- + Can single-step with a debugger.
- + No process or thread control overhead.
  - Design of choice for high-performance Web servers and search engines.
- Significantly more complex to code than process- or thread-based designs.
- Can be vulnerable to denial of service attacks
  - How?

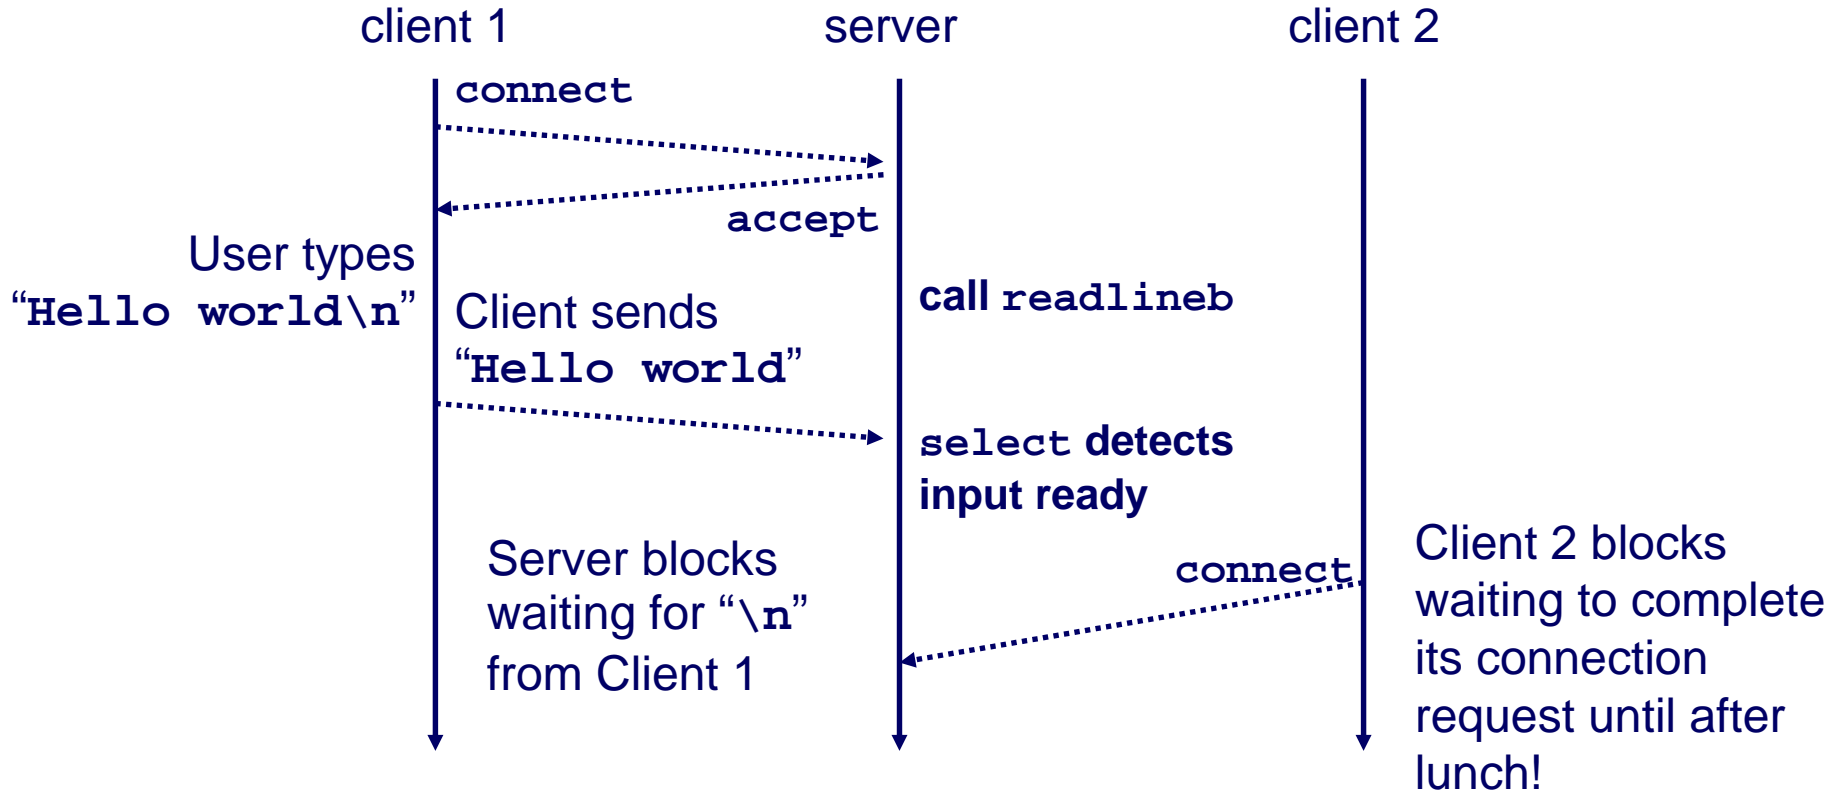
# Attack #1

## Overwhelm Server with Connections

- Limited to `FD_SETSIZE` – 4 (typically 1020) connections

## Defenses?

# Attack #2: Partial Lines



- Client gets attention of server by sending partial line
- Server blocks until line completed

# Flaky Client

```
while (Fgets(buf, MAXLINE, stdin) != NULL) {  
    Rio_writen(clientfd, buf, strlen(buf)-1);  
    Fgets(buf, MAXLINE, stdin); /* Read & ignore line */  
    Rio_writen(clientfd, "\n", 1);  
    Rio_readlineb(&rio, buf, MAXLINE);  
    Fputs(buf, stdout);  
}
```

- Sends everything up to newline
- Doesn't send newline until user types another line
- Meanwhile, server will block

# Implementing a Robust Server

## Break Up Reading Line into Multiple Partial Reads

- Every time connection selected, read as much as is available
- Construct line in separate buffer for each connection

## Must Use Unix Read

`read(int fd, void *buf, size_t maxn)`

- Read as many bytes as are available from file `fd` into buffer `buf`.
- Up to maximum of `maxn` bytes

## Cannot Use RIO Version

`rio_readn(int fd, void *buf, size_t n)`

- Read `n` bytes into buffer `buf`.
- Blocks until all `n` read or EOF

# Robust Server

```
/*
 * echoserverub.c - A robust, concurrent echo server based on select
 */
#include "csapp.h"

typedef struct { /* represents a pool of connected descriptors */
    int maxfd;          /* largest descriptor in read_set */
    fd_set read_set;    /* set of all active descriptors */
    fd_set ready_set;   /* subset of descriptors ready for reading */
    int nready;         /* number of ready descriptors from select */
    int maxi;           /* highwater index into client array */
    int clientfd[FD_SETSIZE]; /* set of active descriptors */
    char clientbuf[FD_SETSIZE][MAXBUF]; /* set of read buffers */
    int clientcnt[FD_SETSIZE]; /* Count of characters in buffers */
} pool;

int byte_cnt = 0; /* counts total bytes received by server */
```

# Robust Server Loop

```
void check_clients(pool *p)
{
    int i, connfd, n;
    for (i = 0; (i <= p->maxi) && (p->nready > 0); i++) {
        connfd = p->clientfd[i];
        char *buf = p->clientbuf[i]; /* Private buffer */
        int cnt = p->clientcnt[i]; /* Number of chars read so far */
        if ((connfd > 0) && (FD_ISSET(connfd, &p->ready_set))) {
            p->nready--;
            if ((n = Read(connfd, buf+cnt, MAXBUF-cnt)) != 0) {
                byte_cnt += n; cnt += n;
                if (buf[cnt-1] == '\n') {
                    Write(connfd, buf, cnt); /* End of line */
                    p->clientcnt[i] = 0;
                } else
                    p->clientcnt[i] = cnt;
            }
        }
    }
}
```

# NOTE (Just to complicate things...)

If a client sends  $x$  bytes of data in one `write()` call, it is NOT guaranteed that all  $x$  bytes will be received in a single `read()` call by the server.

i.e., the following scenario is possible:

- Client writes “Hello world\n” to server
- Server’s `select()` notices & unblocks, server then calls `read()`
- `read()` returns “Hell”
- A subsequent call to `read()` returns “o w”
- A subsequent call to `read()` returns “orld\n”

**Server’s solution:** maintain a buffer for each of your clients, and only process the buffer’s contents when a message has been received in full (note: the type of application determines what a ‘message’ is, and what indicates it has been fully received)



# Conceptual Model

## Maintain State Machine for Each Connection

- First Version: State is just identity of connfd
- Second Version: State includes partial line + count of characters

## select Determines Which State Machine to Update

- First Version: Process entire line
- Second Version: Process as much of line as is available

## Design Issue

- Must set granularity of state machine to avoid server blocking

# For More Information

W. Richard Stevens, *Unix Network Programming: Networking APIs: Sockets and XTI*, Volume 1, Second Edition, Prentice Hall, 1998.

- THE network programming bible.

Complete versions of original echo client and server are developed in *Computer Systems: A Programmer's Perspective*.

- Available from [csapp.cs.cmu.edu](http://csapp.cs.cmu.edu)
- Compile and run them for yourselves to see how they work.
- Feel free to borrow any of this code.
- But be careful--it isn't sufficiently robust for our programming assignments
  - » Most routines exit when any kind of error encountered

# For More Information

## What's inside the RIO wrappers

```
int Select(int n, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout) {
    int rc;
    if ((rc = select(n, readfds, writefds, exceptfds, timeout)) < 0)
        unix_error("Select error");
    return rc;
}
```

# GDB and Version Control (RCS/CVS)

## Tools to help make programming task simpler

- GDB helps the debugging task
- RCS or CVS help with the task of maintaining your code across multiple revisions

## Neither system is magic

- Debugging
  - Program defensively – check error returns, buffer sizes
  - Build program in a modular fashion
  - Print sensible error messages
    - » May want to do better than the built-in wrapper functions
- Version control
  - Keep a clear idea of who is doing what
  - Build program in a modular fashion
    - » Define interfaces between modules early and try not to change them too much

# Debugging with GDB

## Prepare program for debugging

- Compile with “-g” (keep full symbol table)
- Don't use compiler optimization (“-O”, “-O2”, ...)

## Two main ways to run gdb

- On program directly
  - `gdb progname`
  - Once gdb is executing we can execute the program with:
    - » `run args`
    - » Can use shell-style redirection e.g. `run < infile > /dev/null`
- On a core (post-mortem)
  - `gdb progname core`
  - Useful for examining program state at the point of crash

## Extensive in-program documentation exists

- `help` (or `help <topic>` or `help <command>` )

# Controlling Your Program With GDB

## Stopping your program with breakpoints

- Program will run until it encounters a breakpoint
  - To start running again: `cont`
- Break command format
  - `break foo.c:4` stops at the 4<sup>th</sup> source line of `foo.c`
  - `break 16` stops at the 16<sup>th</sup> source line of the current source file
  - `break main` stops upon entry to the `main()` function

## Stop your program with SIGINT (CTRL-C)

- Useful if program hangs (sometimes)

## Stepping through your program

- `step N` command: steps through N source lines (default 1)
- `next` is like `step` but it treats function calls as a single line

## Hint: avoid writing mega-expressions

- Hard to step through `foo(bar(tmp = baz(), tmp2 = baz2()))`

# Examining the State of Your Program

## `backtrace ( bt for short)`

- Shows stack trace (navigate procedures using `up` and `down`)
- `bt full` prints out all available local variables as well

## `print EXP`

- Print the value of expression
- Can print out values in variables

## `x/<count><format><size> ADDR`

- Examine a memory region at `ADDR`
- Count is the number of items to display (default: 1)
- Format is a single letter code
  - `o`(octal), `x`(hex), `d`(decimal), `u`(unsigned decimal), `t`(binary), `f`(float), `a`(address), `i`(instruction), `c`(char) and `s`(string)
- Size is the size of the items (single letter code)
  - `b`(byte), `h`(halfword), `w`(word), `g`(giant, 8 bytes)

# Version Control with RCS

## Version control systems:

- Maintain multiple versions of a file
  - Allow rollback to old versions
  - Enforce documentation of changes
- Allows multiple programmers to work on a project without accidentally editing the same file
  - Files must be 'checked out' for reading or writing

## RCS maintains a database of all revisions

- Make a subdirectory called 'RCS' in each working directory
  - Otherwise RCS will do its business in your directory – ugly!
- If your file is called 'assignment1/foo.c', RCS keeps update history in 'assignment1/RCS/foo.c,v'
- Current version of 'foo.c' is maintained in 'foo.c,v'
  - 'deltas' allow retrieval of older versions



# Creating RCS Files

After making 'RCS' subdirectory...

Initialize RCS for your file `mysource.c` (assume you have already created it) by checking it in ( `ci` )

```
[geoffl@ux3 ~/tmp]$ ci mysource.c
RCS/mysource.c,v  <--  mysource.c
enter description, terminated with single '.' or end of file:
NOTE: This is NOT the log message!
>> This source file contains a simple algorithm to solve the Halting
Problem.
>> .
initial revision: 1.1
done
```

Can also create a blank file with `rscs -i mysource.c`

Either way, this produces version 1.1

# Checking out files

In previous example, after `ci`, `mysource.c` is gone!

To retrieve `mysource.c`, use `co` command ('check out')

```
[geoffl@ux3 ~/tmp]$ ls -l
total 2
drwxr-xr-x    2 geoffl    users          2048 Jan 19 15:39 RCS
[geoffl@ux3 ~/tmp]$ co mysource.c
RCS/mysource.c,v --> mysource.c
revision 1.1
done
[geoffl@ux3 ~/tmp]$ ls -l
total 3
-r--r--r--    1 geoffl    users          137 Jan 19 15:39 mysource.c
drwxr-xr-x    2 geoffl    users          2048 Jan 19 15:39 RCS
```

**Note:** permissions don't let us change `mysource.c`

To change `mysource.c`, must acquire lock

- `co -l mysource.c` locks `mysource.c` so no one else can change it
- Use `ci` to check the code back in when done (adding a log message)

# Versions

Each version of the file has a version number

- “release.revision” format – e.g. 4.2 is release 4, revision 2
- Doesn’t necessarily correspond to anything about real world version numbers

By default, each `ci` of a changed file increments revision number by 1

Can use `-r` flag to specify version numbers

- Use this with `co` to retrieve old versions
- Use this with `ci` to specify what a new version should be called
  - Note: can’t go backwards!
  - `ci -r1.8 mysource.c` will check in `mysource.c` with version number 1.8
  - `ci -r2 mysource.c` will check in `mysource.c` with version 2.1

# Version Control with CVS

**Similar to RCS, but newer and with more functionality**

- Like RCS, Maintains multiple versions of a file
  - Allow rollback to old versions
  - Enforce documentation of changes
- Allows multiple programmers to work on the same file at the same time
  - Upon checkin, if conflicts exist, the user is notified and can resolve them manually using a diff program (i.e. diff on UNIX, ExamDiff on Windows)
- Easy to use on the UNIX command line
- Has a very, very friendly client for Windows users that integrates into your folder's file management system called "TortoiseCVS" <http://www.tortoisecvs.org/>

# Version Control continued

## RCS and CVS both maintain a database of all revisions

- Make a subdirectory called 'RCS' in each working directory
  - Otherwise RCS will do its business in your directory – ugly!
  - Note: CVS does this for you automatically
- If your file is called 'assignment1/foo.c', RCS keeps update history in 'assignment1/RCS/foo.c,v', CVS similarly
- Current version of 'foo.c' is maintained in 'foo.c,v'
  - 'deltas' allow retrieval of older versions

## More information about CVS (my personal choice)

- <http://www.nongnu.org/cvs/>
- <http://www.tortoisecvs.org/> for Windows users

# More information...

## GDB

- Official GDB homepage:  
<http://www.gnu.org/software/gdb/gdb.html>
- GDB primer: <http://www.cs.pitt.edu/~mosse/gdb-note.html>

## RCS

- Look at `man rcs`, `man rcsintro`
- Official RCS homepage:
  - <http://www.cs.purdue.edu/homes/trinkle/RCS/>
- Other useful features
  - `ci -l`: check-in a version but keep the file and the lock
  - `ci -u`: check-in a version but keep a read-only version of file
  - `rcsdiff`: display differences between versions
  - `rcsmerge`: merge changes in different versions of a file
  - Note: you can break locks if necessary
    - » RCS will send e-mail to owner of broken lock

# Tips for the past

**Find a partner that doesn't procrastinate.**

**Schedule a \*daily\* meeting time. DAILY.**

- Do a couple hours of work each day.
- Even meeting 3 days a week for a few hours, my partner and I pulled multiple all-nighters. Avoid this by meeting daily.
- Your implementations for each project can be expected to be 5,000 lines of code, plus or minus a few thousand (ours were between 4,000-6,000 lines). Divide that by days and it's not as daunting.

**START THE DAY YOU RECEIVE THE PROJECT.**

- My personal impression from last year is that the majority of failed/mostly unsuccessful projects failed due to time-related pressure, not content/understanding material pressure
- If you have time to do it, you can do it well. If not...





# More tips for the past

If you want to work directly on windows, you can use MinGW or Cygwin, **\*BUT TEST ON UNIX\***

- <http://www.mingw.org/>
- <http://www.cygwin.com/>

Most of the APIs you will use will have many functions that return -1 to indicate error. Unlike previous classes, you must be able to recover when appropriate, and handle these errors on a case-by-case basis.

Make buffers for **\*each\*** of your clients; don't process a command until you know you've received the entire thing.

# Finally...

**These projects are about implementing specifications.**

**Read the writeup at least twice before starting.**

**PLAN on paper (even draw) BEFORE coding.**

**PLAN, PLAN, PLAN.**

**Then coding and debugging will be much easier, and  
you won't have to delete hundreds of lines of useless  
code over and over again.**

**Good luck!**